

Aurélien Rainone
Go Software Engineer

How to Instrument and Monitor a Go Application

The Student Hotel

11th December 2019



- Is my program slow, can it go faster?
- How does my app perform in production?
- How does it perform under load?

- Monitor a CLI application
 - Tracing
 - Profiling
- Monitor a HTTP server
 - Time Series Database
 - Metrics
 - Prometheus

Instrument a command line application

Mandelbrot Set



```
$ git clone github.com/ar1/monitoring  
$ cd monitoring  
$ cd fractal
```

```
for each frame:
    for each pixel:
        // c = x + yi
        c = complex(pixel.x, pixel.y)
        z = complex(0, 0)

        // does z 'escapes' to infinity
        for n=0 => max_iteration:
            z = z2+c
            if mod(z) > some_number:
                break

        // use n to color the pixel
        if n == max_iterations:
            pixel.color = black
        else:
            pixel.color = f(n)
```

https://en.wikipedia.org/wiki/Mandelbrot_set

Execution Tracer

golang.org/pkg/runtime/trace/
go tool trace

- To spot concurrency problems
- See the interaction between our program and the Go runtime (scheduler, GC) and ultimately with the OS

Benchmarking

golang.org/pkg/testing/

go test -bench ...

golang.org/x/perf/cmd/benchstat

- To measure performance at the function level
- Compare code performance

Profiling

golang.org/pkg/runtime/pprof/
go tool pprof -http=:

CPU, Memory, Contention, etc.

- Using the pprof package directly (Start/Stop)
- Adding a profiling HTTP handler to our program
- go test -cpuprofile or -memprofile (only CPU and Memory)

CPU Profiling

```
go test -bench . -run $^ -cpuprofile cpu.out
```

Stops your program every 10ms and records the stack trace of each goroutine.

Memory Profiling

```
go test -bench . -run $^ -memprofile cpu.out
```

Records the stack trace at each heap allocation.

alloc_objects: objects allocated

alloc_space: bytes allocated


inuse_space: objects allocated and not GC'ed

inuse_objects: bytes allocated and not GC'ed



Service Level Metrics

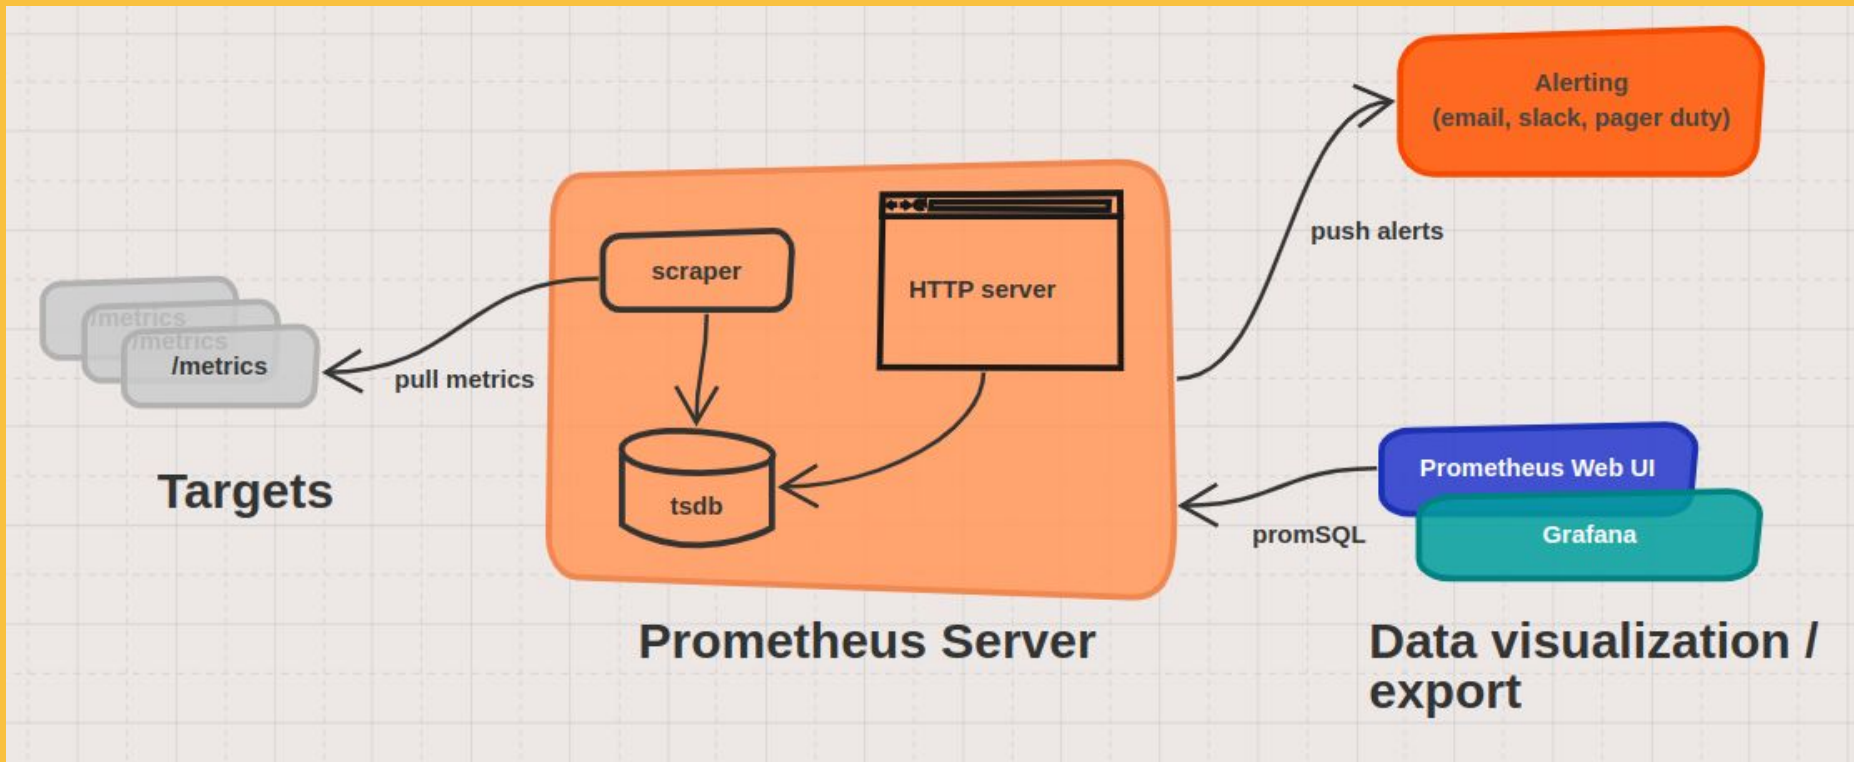
Time Series Databases



Timestamp	Metric ID	Value
13:01:45	12	10327
13:01:47	12	10500
13:01:48	14	3.14159275



Prometheus




Metric types

Counter


- can only increase
- number of requests
- number of errors
- ...
- compute rates

Counter


$$rate = \frac{N_{t1} - N_{t0}}{t_1 - t_0}$$



Gauge




- can increase and decrease
- temperature
- CPU load
- memory usage
- ...




- value distribution

Histogram




- sample values
- counts them in bucket

Histogram



- median
- percentiles

Histogram



Hacking

```
$ git clone github.com/ar1/monitoring  
$ cd monitoring  
$ cd cache
```

- **Add** a key-value pair to the cache:

`http://localhost:8080/add?k=key&v=value`

- **Get** a cached value

`http://localhost:8080/get?k=key`

- Count all requests
- Plot the number of requests per minute

TIPS:

use an http middleware
use rate (promQL)

**Plot cache
misses/hits**

- Record request duration
- Plot the median (50th percentile) over the last 10m

TIPS:

modify the previous middleware
use microseconds as unit
use histogram_quantile (promQL)

- Add a label named “endpoint” on request duration to have separate values for “add” and “get”

TIPS:

you can still do it with a middleware

use `promauto.NewHistogramVec`

Thank you!

Aurélien Rainone

aurelien@develer.com

The logo for 'develer' features the word in a white, lowercase, sans-serif font. A horizontal orange bar is positioned above the 'e'.

www.develer.com